

Genome assembly and annotation

GastroPak Microbial Bioinformatics Workshop

DR CHRISTOPHER QUINCE

Earlham/Quadram Group Leader




Overview

- Brief introduction to microbial genome assembly
- Assembly evaluation
- Microbial genome annotation

Obtaining samples for genome sequencing

- Microbial genome sequencing usually proceeds from cultured isolates
- Streak or spread plate method
- We typically know what species should be present as a single homogenous strain
- Good practice to check with Sanger sequencing of 16S rRNA
- Need to consider DNA extraction protocol but typically can obtain high quality DNA
- Can also perform 'sweep metagenomics' by sequencing whole plate prior to isolation
- Multiple strains same species



Single-cell microbial genome sequencing

- FACs sorting of individual cells
- Sequencing following multiple displacement amplification (MDA)

The trajectory of microbial single-cell sequencing

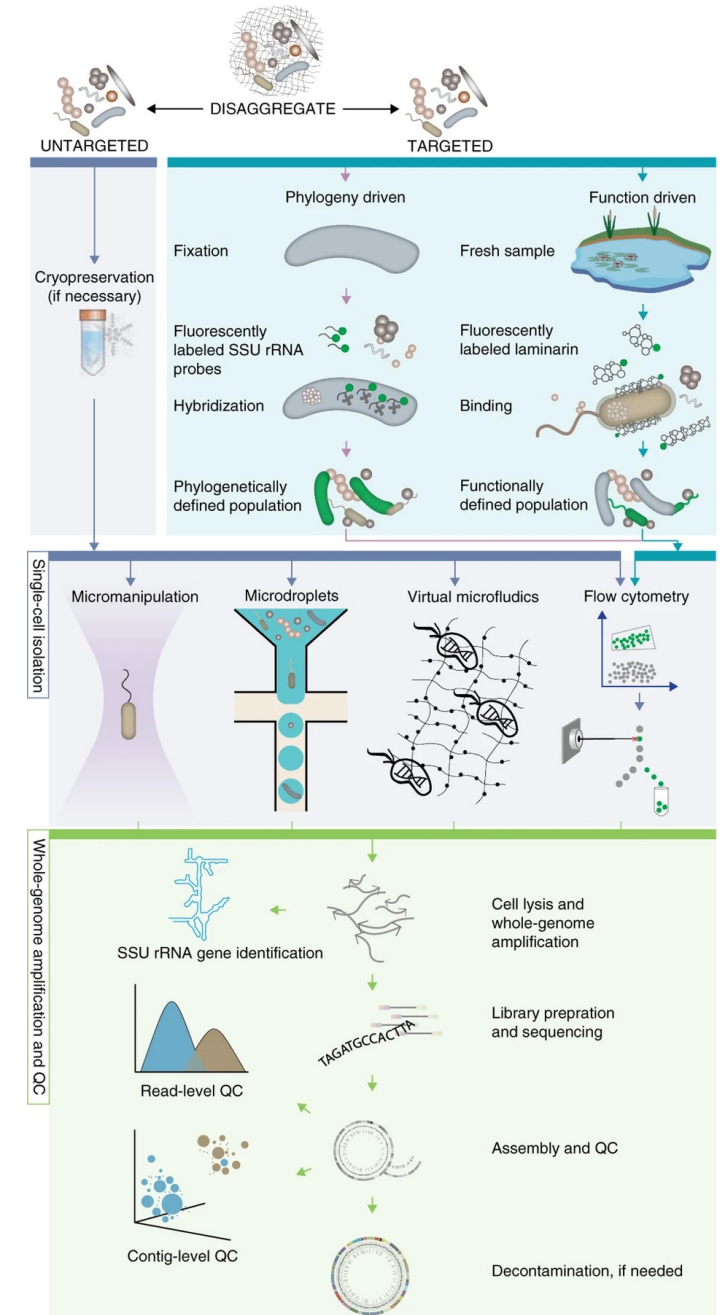
[Tanja Woyke](#) , [Devin F R Doud](#) & [Frederik Schulz](#)

[Nature Methods](#) **14**, 1045–1054 (2017) | [Cite this article](#)

9091 Accesses | 90 Citations | 57 Altmetric | [Metrics](#)


<https://www.nature.com/articles/nmeth.4469/>

www.earlham.ac.uk



The genome assembly problem

Shotgun genomics DNA is physically broken and replicated into millions of small fragments



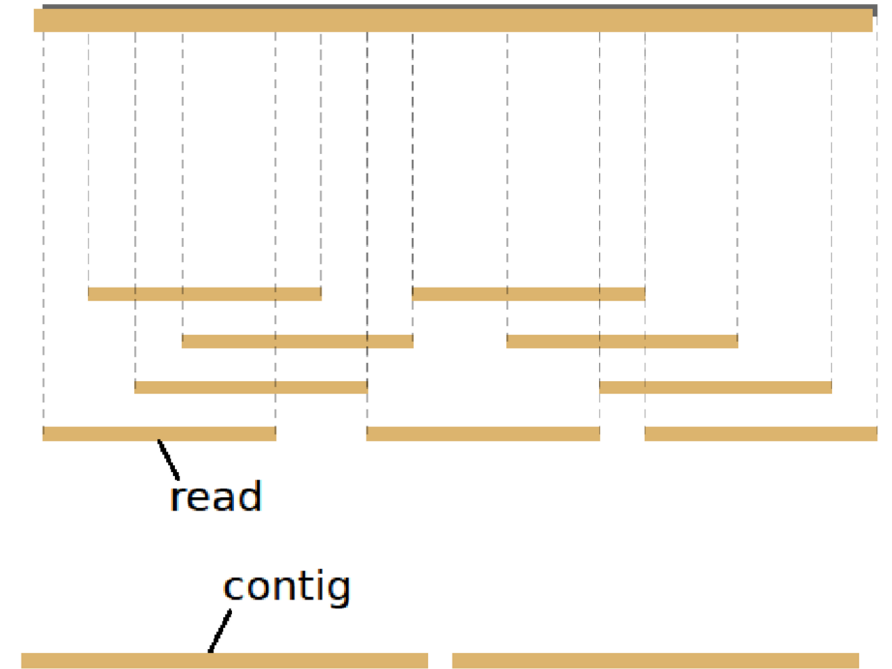
Assembly methods

- Overlap-Layout-Consensus
 - Overlap/String graph
- de Bruijn graph

genome
(unknown)

sequenced reads:

overlapping
sub-sequences,
covering
the genome
redundantly



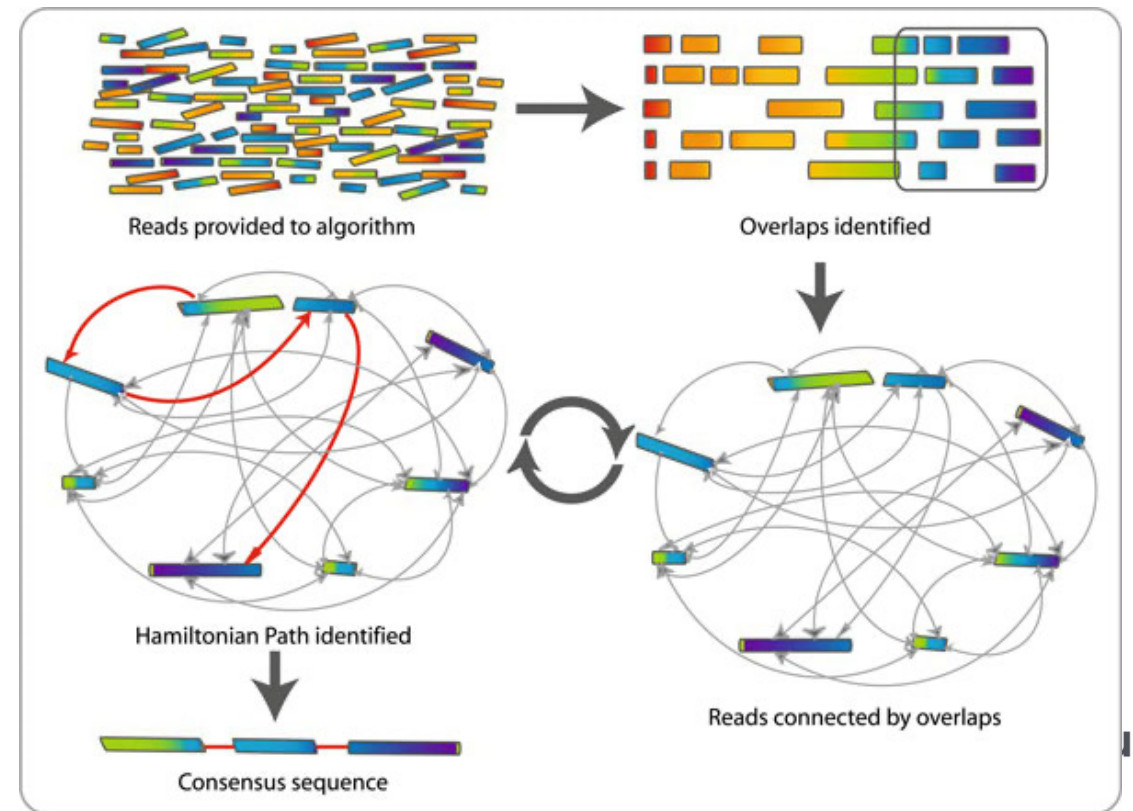
assembly
hypothesis of
the genome

Short-read metagenomics assembly

- Assembly consists of looking for overlaps between short reads

Read 1: ACTCGTCGTG → ACTCGTCGTG
Read 2: GTGCGATTCC → GTGCGATTCC → Contig: ACTCGTCGTGCGATTCC

- Two methods for detecting read overlaps, Overlap Layout Consensus (OLC) and de Bruijn graph assemblers
- OLC is more intuitive:
 - Overlap: reads are compared to one another to identify overlapping regions (MinHash)
 - Layout: reads are positioned based on overlaps
 - Create graph with nodes (reads) and edges (overlaps)
 - Hamiltonian path visits every vertex once
 - Consensus: multiple sequence alignment is carried out to produce a consensus sequence
- Scales badly with read number – why?



Detecting sequence overlaps

- Requires pairwise comparisons scales as N^2
- Comparison requires alignment of two reads
- Why is alignment needed?
- This is computationally expensive (see below)
- Speed up with minimizer methods

ACTGGGCAATGACTC
ATTGCTCACCATCCG

ACTGGGC**AATGACTC**
 ATTG-CTCACCATCCG

Edit distance two
One overlap one substitution
Overlap 8 bp

K-mers

- k-mers are substrings of length k contained within a biological sequence e.g. trimers
- Any sequence of length L will contain $L - k + 1$ *k-mers*
- How many unique k=4 tetramers are there?

Sequence: ATCGATCAC

3-mer #0: ATC

3-mer #1: TCG

3-mer #2: CGA

3-mer #3: GAT

3-mer #4: ATC

3-mer #5: TCA

3-mer #6: CAC

<https://bioinfologics.github.io/post/2018/09/17/k-mer-counting-part-i-introduction/> Bernardo Clavijo

K-mers (continued)

- Often convenient to consider ‘canonical’ kmers where a kmer and its reverse complement are equivalent e.g. ACTG == CAGT
- This is because a DNA sample and hence sequenced reads (almost) always contains sequence from both strands
- How many k = 3 and k = 4 canonical kmers are there?
- Oten use odd values of k to avoid palindromic k-mers
- Why are k-mers useful:
 - Enable efficient sequence comparison without alignment
 - Scalable short read assembly

Offset	0	1	2	3	4	5	6
3-mer	ATC	TCG	CGA	GAT	ATC	TCA	CAC
Reverse Complement	GAT	CGA	TCG	ATC	GAT	TGA	GTG
Canonical	ATC	CGA	CGA	ATC	ATC	TCA	CAC

Sequence: ATCGATCAC
3-mer #0: ATC
3-mer #1: TCG
3-mer #2: CGA
3-mer #3: GAT
3-mer #4: ATC
3-mer #5: TCA
3-mer #6: CAC

Canonical k-mer	Total count
ATC (incl. GAT)	3
CGA (incl. TCG)	2
TCA	1
CAC	1
Sum	7

de Bruijn graph assembly

Steps:

1. Identify kmers in reads
2. Take all (k-1)-mers from the set of k-mers, e.g. ATG, TGC → AT, TG, GC
3. Construct a multi-graph with nodes being k-1-mers; draw an edge between two k-1 mers only if the two k-1 mers are taken from the same read. **Ex:** GCT & CTA
4. Graph constructed this way is guaranteed to have a Eulerian trail, follow the trail and connect the nodes to form our original sequence

- Efficient storage space
 - fixed word length
 - Independent of sequencing depth
- How to choose k? → iterative kmers

www.earlham.ac.uk



Uditha Maduranga


1.1K Followers

Computer Science and Engineering Graduate | AWS
Community Builder | Former Intern — Software
Engineer at Sysco LABS | maduranga95@gmail.com

Genome: ATGCTAGCAC

Reads: ATGCTA
GCTAGC
TAGCAC
GCACAT
ACATGC

k-mers: ATG TGC GCT
CTA TAG
AGC GCA CAC
ACA CAT



Earlham Institute


<https://medium.com/towards-data-science/genome-assembly-using-de-bruijn-graphs-69570efcc270>

Short read assemblers

- 1) de Bruijn **graph** construction



- 2) Likely **sequencing errors** are removed.



- 3) **Variations** (e.g. SNPs, similar repetitions) are removed.

→ **Collapses strains**


- 4) **Simple paths** (i.e. contigs) are returned.



- 5) Extra steps: repeat-resolving, scaffolding

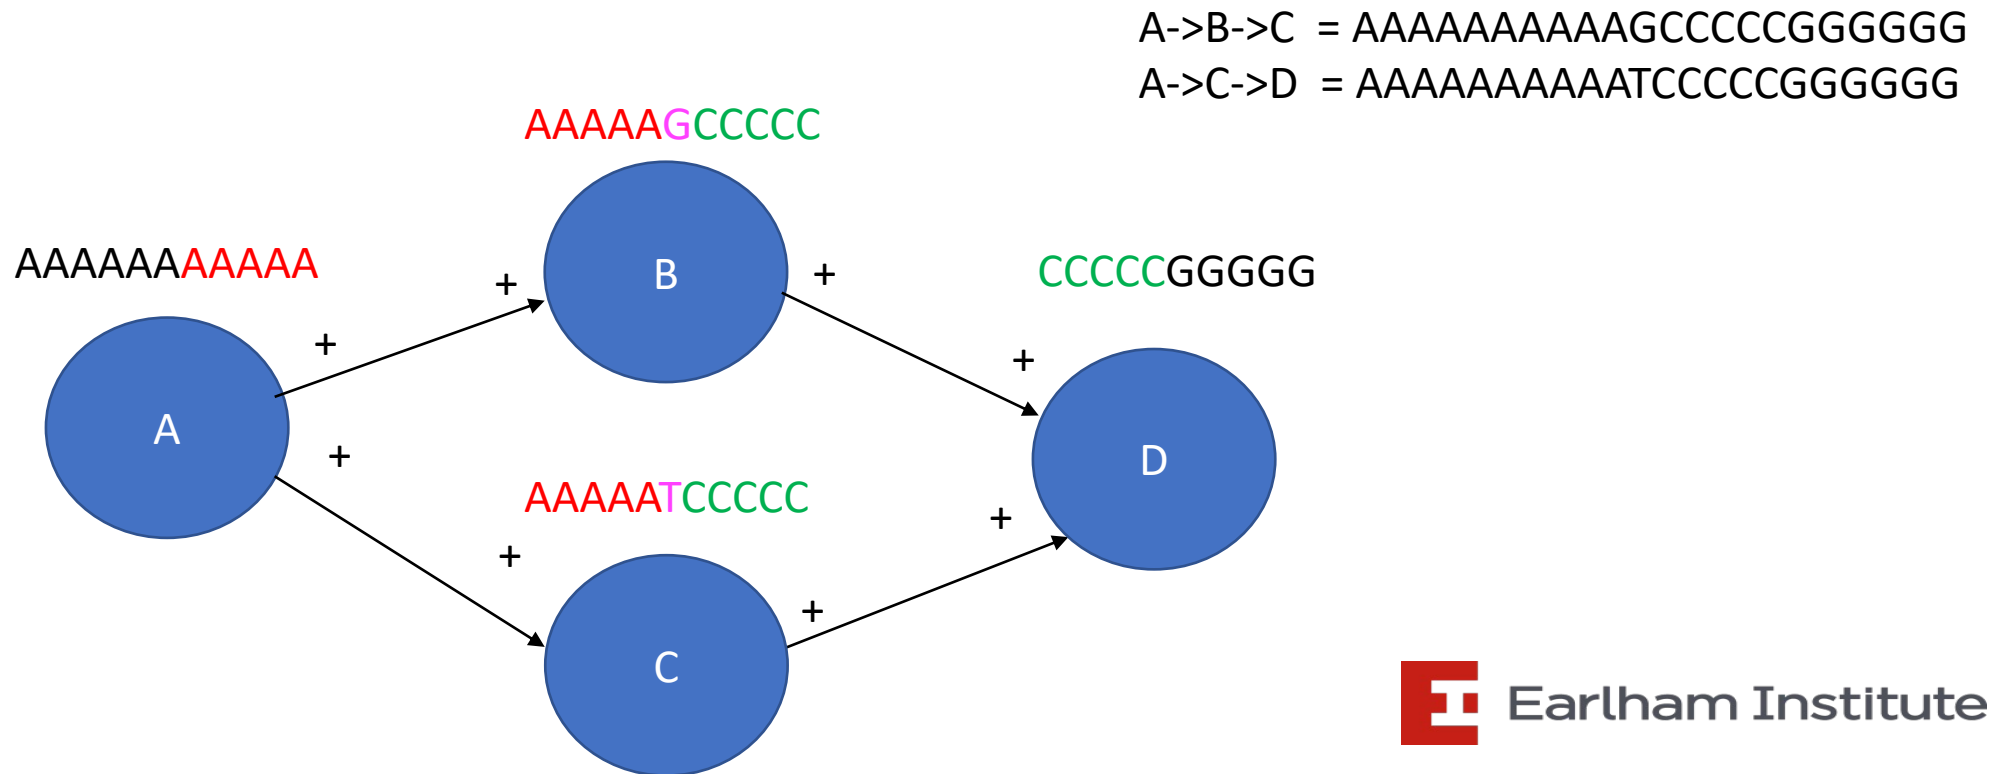
Genome repeats

- Difficulties arise because of 'repeat' regions on genomes that exceed read length



Assembly graph represents uncertainty


- Produced by all assemblers following compaction (de Bruijn graphs) or removal of transitive reads (string graphs)
- Nodes represent sequence (unitigs) and edges overlaps
- Represents fundamental uncertainty in possible paths and hence sequences



Assembly evaluation

- Evaluate an assembly by looking at the following statistics:
 - The size of the assembly: sum of contig lengths
 - N50: Given a set of contigs, the N50 is defined as the sequence length of the shortest contig at 50% of the total assembly length
 - Example: consider 9 contigs with lengths 2,3,4,5,6,7,8,9,10; their sum is 54, half of the sum is 27, and the size of the genome also happens to be 54. 50% of this assembly would be $10 + 9 + 8 = 27$ (half the length of the sequence): N50=8
- What factors effect assembly quality? Key one is average depth of coverage:

Depth of coverage




Average depth of coverage of genome

$$\langle c \rangle = \frac{RN}{L}$$

- R is read length
- N is number of reads
- L is genome length

Breadth of coverage

- At low average coverage depths then genome is fragmented through chance
- Bread of coverage is fraction of genome covered by at least one read or equivalent fraction of bases with $c_n > 0$



$$b = \frac{\sum_{n=1}^L c_n = 0}{L}$$

Lander-Waterman statistics

- Poisson statistics allow us to estimate breadth of coverage from average coverage depth:

$$b = 1 - \exp(-\langle c \rangle)$$

- At an average coverage of 5 then 99.3% of genome will be sequenced
- The expected number of contigs:

$$g = N \exp(-\langle c \rangle)$$

- Coverage depth of 5 with 300 bp reads and 5 Mbp genome what is the expected number of contigs?
- This is an best case prediction due to repeats


Genome assembly software

- For short reads typically use iterative k-mer de Bruijn graph assembler:
 - SPAdes: <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3342519/>
 - Megahit
- For long reads OLC:
 - Nanopore: Flye or Canu
 - HiFi PacBio: hifiasm

Contamination

- Potential issue is presence of sequence from organisms other than the target genome
- Can explore this very simply by looking at k-mer frequency spectrum which for bacteria should be unimodal
- Slightly more sophisticated blob plots

KAT: a K-mer analysis toolkit to quality control NGS datasets and genome assemblies

Daniel Mapleson, Gonzalo Garcia Accinelli, George Kettleborough, Jonathan Wright, Bernardo J Clavijo 

Bioinformatics, Volume 33, Issue 4, February 2017, Pages 574–576,

<https://doi.org/10.1093/bioinformatics/btw663>

Published: 28 November 2016 **Article history** ▼



PDF

Split View



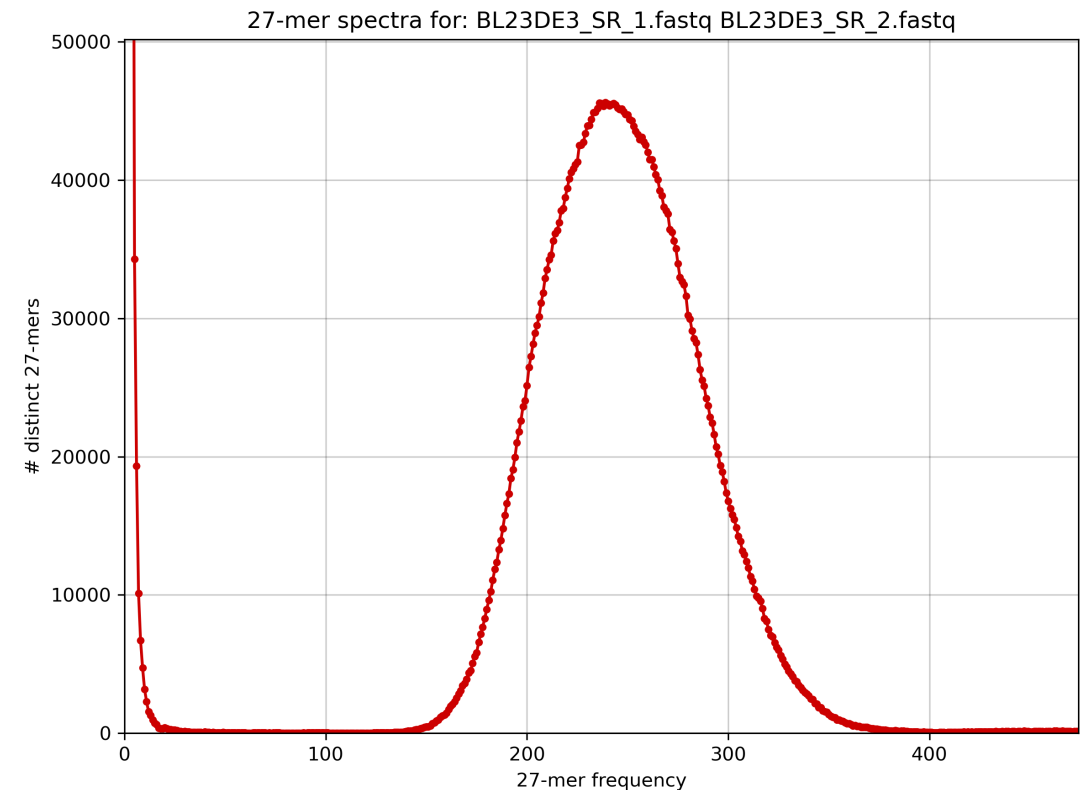
Cite




Permissions



Share ▼





G3 (Bethesda), 2020 Apr; 10(4): 1361–1374.

Published online 2020 Feb 18. doi: [10.1534/g3.119.400908](https://doi.org/10.1534/g3.119.400908)

PMCID: PMC7144090

PMID: [32071071](https://pubmed.ncbi.nlm.nih.gov/32071071/)

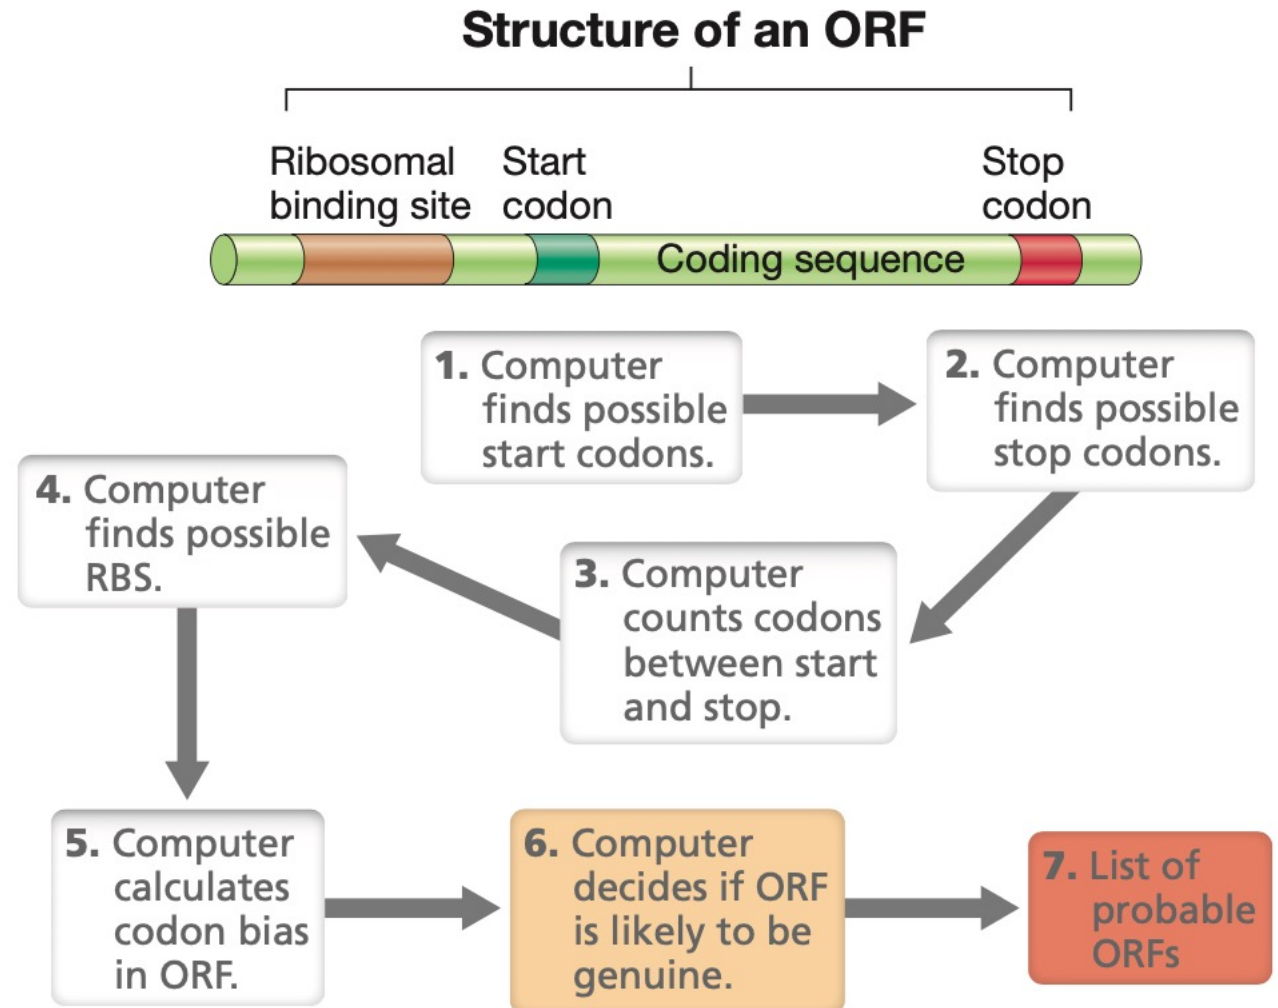
BlobToolKit – Interactive Quality Assessment of Genome Assemblies

Richard Challis,^{†,1} Edward Richards,[‡] Jeena Rajan,[‡] Guy Cochrane,[‡] and Mark Blaxter[†]

► Author information ► Article notes ► Copyright and License information ► [Disclaimer](#)

Genome annotation – Identifying ORFs

- Steps in bacterial genome annotation:
 - Identify open reading frames (ORFs) in contigs
- Typical program for this is Prodigal
- Does require the genetic code to be known
- Prodigal will try and guess code but can get this wrong for exotic microbes



Annotating protein coding ORFs to functions

- Annotating ORFs to potential functions is done by searching against databases
- Essentially based on 'sequence homology'
- We assume that similar sequences perform similar functions
- Search can be performed either with amino acids or nucleotides
 - Pros and cons?
- Complication is need to align sequences


Reference database example

- Database is a set of sequences with labels (annotations)
- Labels can be any feature of interest e.g. Clusters of orthologous genes (COGs)
<http://www.ncbi.nlm.nih.gov/COG/>
- Each COG consists of individual orthologous genes, share common ancestor through speciation, with a particular function:
 - e.g. COG0001 Glutamate-1-semialdehyde aminotransferase
- Current COG based on complete genomes of 1187 bacteria and 122 archaea
- Database consists of 4,877 COGs that include 3,236,575 unique genes mapped to 3,455,867 genomic loci, which represents 73.5% of the 4,401,819 proteins (4,110,746 bacterial and 291,073 archaeal proteins) encoded by the covered species
- Return to databases in later metagenomics lectures

COG database update: focus on microbial diversity, model organisms, and widespread pathogens

[Michael Y Galperin](#), [Yuri I Wolf](#), [Kira S Makarova](#), [Roberto Vera Alvarez](#), [David Landsman](#), and [Eugene V Koonin](#)

[▶ Author information](#) ▶ [Article notes](#) ▶ [Copyright and License information](#) [Disclaimer](#)



Functional classification of prokaryotic (COGs) and eukaryotic (KOGs) clusters of orthologs. Designations of functional categories: A, RNA processing and modification (not used for prokaryotic COGs), B, chromatin structure and dynamics, C, energy production and conversion, D, cell cycle control and mitosis, E, amino acid metabolism and transport, F, nucleotide metabolism and transport, G, carbohydrate metabolism and transport, H, coenzyme metabolism, I, lipid metabolism, J, translation, K, transcription, L, replication and repair, M, cell wall/membrane/envelope biogenesis, N, Cell motility, O, post-translational modification, protein turnover, chaperone functions, P, Inorganic ion transport and metabolism, Q, secondary metabolites biosynthesis, transport and catabolism, T, signal transduction, U, intracellular trafficking and secretion, Y, nuclear structure (not applicable to prokaryotic COGs), Z, cytoskeleton (not applicable to prokaryotic COGs); R, general functional prediction only (typically, prediction of biochemical activity), S, function unknown. The numbers were obtained after subtracting the COGs that consisted entirely of proteins from unicellular eukaryotes from the COG collection.

Sequence alignment

- Alignment can be:
 - Global
 - Local
- Global matches whole sequence
 - search ORF against a database of ORFs
- Local matches part of query to reference database
 - search ORF against a genome

Local Alignment

Target Sequence	5' ACTACTAGATTACTTACGGATCAGGTACTTTAGAGGCTTGCAACCA 3'
Query Sequence	5' TACTCACGGATGAGGTACTTTAGAGGC 3'

Global Alignment

Target Sequence	5' ACTACTAGATTACTTACGGATCAGGTACTTTAGAGGCTTGCAACCA 3'
	5' ACTACTAGATT----ACGGATC--GTACTTTAGAGGCTAGCAACCA 3'
Query Sequence	

Sequence alignment (cont.)

- Sequence alignment can be performed exactly given a scoring system:
 - E.g. Match (+1), mismatch (-1) or indel (-2)
- Needleman-Wunsch (Smith-Waterman) algorithm finds global (local) optimal alignment for any scoring system
- These are forms of dynamic programming and scale with L^2 where L is the sequence length
- Simple scoring suitable for nucleotides
- Amino acids have different properties and frequencies motivates complex scoring matrices e.g. BLOSUM62 and PAM

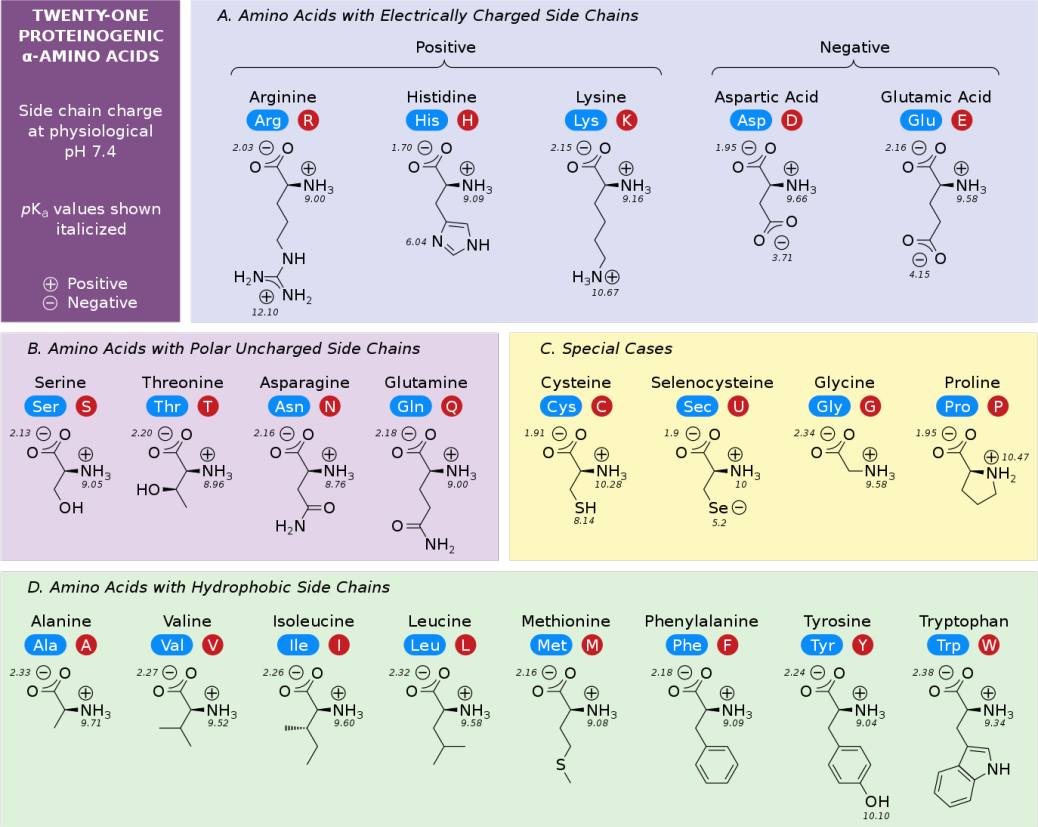
Needleman-Wunsch

match = 1mismatch = -1gap = -1

		G	C	A	T	G	C	G	
		0	-1	-2	-3	-4	-5	-6	-7
G	-1	1	0	-1	-2	-3	-4	-5	
A	-2	0	0	1	0	-1	-2	-3	
T	-3	-1	-1	0	2	1	0	-1	
T	-4	-2	-2	-1	1	1	0	-1	
A	-5	-3	-3	-1	0	0	0	-1	
C	-6	-4	-2	-2	-1	-1	1	0	
A	-7	-5	-3	-1	-2	-2	0	0	

The diagram illustrates the Needleman-Wunsch algorithm for sequence alignment. The matrix shows scores for all subproblems. The optimal alignment path is highlighted with blue arrows, starting from the top-left cell (0,0) and ending at the bottom-right cell (8,8). The path follows the sequence GGCAGCTG aligned with TTTACGAA. Red arrows indicate alternative paths or gaps.

BLOSUM: Blocks Substitution Matrix



	C	S	T	A	G	P	D	E	Q	N	H	R	K	M	I	L	V	W	Y	F	
C	9																				C
S	-1	4																			S
T	-1	1	5																		T
A	0	1	0	4																	A
G	-3	0	-2	0	6																G
P	-3	-1	-1	-1	-2	7															P
D	-3	0	-1	-2	-1	-1	6														D
E	-4	0	-1	-1	-2	-1	2	5													E
Q	-3	0	-1	-1	-2	-1	0	2	5												Q
N	-3	1	0	-2	0	-2	1	0	0	6											N
H	-3	-1	-2	-2	-2	-2	-1	0	0	1	8										H
R	-3	-1	-1	-1	-2	-2	-2	0	1	0	0	5									R
K	-3	0	-1	-1	-2	-1	-1	1	1	0	-1	2	5								K
M	-1	-1	-1	-1	-3	-2	-3	-2	0	-2	-2	-1	-1	5							M
I	-1	-2	-1	-1	-4	-3	-3	-3	-3	-3	-3	-3	-3	1	4						I
L	-1	-2	-1	-1	-4	-3	-4	-3	-2	-3	-3	-2	-2	2	2	4					L
V	-1	-2	0	0	-3	-2	-3	-2	-2	-3	-3	-3	-2	1	3	1	4				V
W	-2	-3	-2	-3	-2	-4	-4	-3	-2	-4	-2	-3	-3	-1	-3	-2	-3	11			W
Y	-2	-2	-2	-2	-3	-3	-3	-2	-1	-2	2	-2	-2	-1	-1	-1	-1	2	7		Y
F	-2	-2	-2	-2	-3	-4	-3	-3	-3	-3	-1	-3	-3	0	0	0	-1	1	3	6	F
	C	S	T	A	G	P	D	E	Q	N	H	R	K	M	I	L	V	W	Y	F	

Practical alignment software

- Practical alignment algorithms use a heuristic to approximate exact alignments
- BLAST (basic **local alignment** search tool) – [Altschul et al. 1990](#)
 - Uses seeds similar to protein k-mers to quickly find matches then extends them
 - Different modes, blastn, blastp, blastx
 - Highly sensitive can find very distant hits but relatively slow
 - Uses BLOSUM matrices
 - Still useful because of website where whole of NCBI is indexed: <https://blast.ncbi.nlm.nih.gov/Blast.cgi>
 - Or when database is small
- DIAMOND ([Buchfink et al. 2015](#))
 - Drop-in replacement for BLAST based on double indexing
 - Similar accuracy but 10,000 times faster
- Mapping software sacrifice sensitivity for speed, accurate but only within certain level of similarity
 - E.g. BWA or Minimap
- Hidden Markov models are sensitive but lack precision for very similar sequences
 - E.g. HMMER

Measures of hit quality used by BLAST

- Percent identity: no. of matches divided by length of alignment
- Bigger database more matches by chance
- Bit-score: Log2 of the database size necessary to find a match this good by chance
 - Larger is better
- E-value: no. of sequences expected to match with this quality when searching against a random database of the same size
 - Smaller is better

$$E = \frac{mn}{2^b}$$

E: e-value

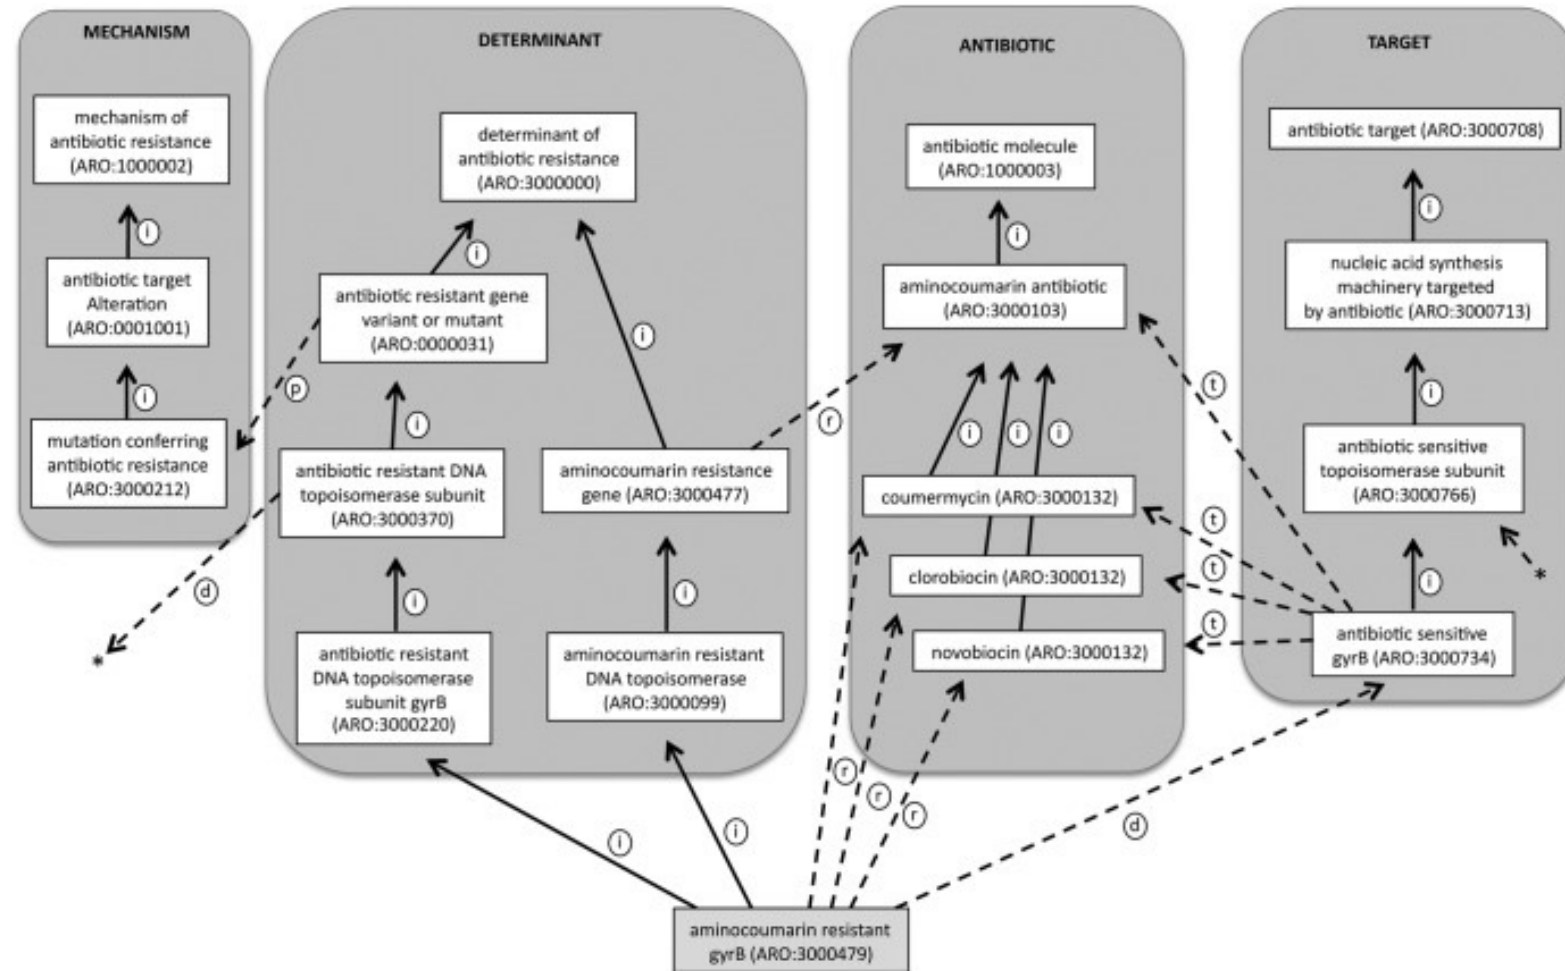
b: bit-score

m: query sequence length

n: database size

Annotation of anti-microbial resistance genes

- Need database of ARGs many examples: e.g. The Comprehensive Antibiotic Resistance Database (CARD)
- CARD (2020 update): 5,148 Reference Sequences mapped to AMR ontology
<https://card.mcmaster.ca>
- Uses BLAST but with custom bit score threshold for each gene



Identifying RNA genes

RNAmmmer: consistent and rapid annotation of ribosomal RNA genes

[Karin Lagesen](#),^{1,2,*} [Peter Hallin](#),³ [Einar Andreas Rødland](#),^{1,2,4,5} [Hans-Henrik Stærfeldt](#),³ [Torbjørn Rognes](#),^{1,2,4} and [David W. Ussery](#)^{1,2,3}

► [Author information](#) ► [Article notes](#) ► [Copyright and License information](#) [Disclaimer](#)

- This can be done using hidden Markov models
- Hidden Markov models are trained with a database and can then with very high sensitivity detect similar sequence
- rRNA sequence is very different from protein coding sequences makes it quite a straightforward problem

Pipelines for genome annotation

- Command line pipelines e.g. Prokka
- Web based servers e.g. RAST <https://rast.nmpdr.org> or IMG <https://img.jgi.doe.gov>

JOURNAL ARTICLE

Prokka: rapid prokaryotic genome annotation FREE

[Torsten Seemann](#) [Author Notes](#)

Bioinformatics, Volume 30, Issue 14, July 2014, Pages 2068–2069,
<https://doi.org/10.1093/bioinformatics/btu153>

Published: 18 March 2014 **Article history** ▼






 PDF  Split View  Cite  Permissions  Share ▼

Table 1. Feature prediction tools used by Prokka

Tool (reference)	Features predicted
Prodigal (Hyatt 2010)	Coding sequence (CDS)
RNAmmer (Lagesen et al. , 2007)	Ribosomal RNA genes (rRNA)
Aragorn (Laslett and Canback, 2004)	Transfer RNA genes
SignalP (Petersen et al. , 2011)	Signal leader peptides
Infernal (Kolbe and Eddy, 2011)	Non-coding RNA

Other things you may want to do with a genome

- Find prophages
- Virulence factors
- Annotate plasmids
- Variant detection
- Pangenomes
- Phylogenetic trees: FastTree, RaxML, IQTree

Future directions

- Functional annotation through structure (AlphaFold)